

Mileage Prediction

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: d=pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv')
d.head()
```

Out[2]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino

```
In [4]: d.nunique()
```

```
Out[4]: mpg          129
cylinders           5
displacement        82
horsepower          93
weight             351
acceleration        95
model_year          13
origin              3
name              305
dtype: int64
```

Data processing

In [5]: d.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg              398 non-null    float64
1   cylinders         398 non-null    int64
2   displacement      398 non-null    float64
3   horsepower        392 non-null    float64
4   weight            398 non-null    int64
5   acceleration      398 non-null    float64
6   model_year        398 non-null    int64
7   origin            398 non-null    object
8   name              398 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

In [6]: d.describe()

Out[6]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
count	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.010050
std	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	3.697627
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	79.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000

In [7]: d.corr()

Out[7]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
mpg	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289	0.579267
cylinders	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419	-0.348746
displacement	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684	-0.370164
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361
weight	-0.831741	0.896017	0.932824	0.864538	1.000000	-0.417457	-0.306564
acceleration	0.420289	-0.505419	-0.543684	-0.689196	-0.417457	1.000000	0.288137
model_year	0.579267	-0.348746	-0.370164	-0.416361	-0.306564	0.288137	1.000000

Missing Values

```
In [16]: d=d.dropna()
```


```
In [17]: d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   mpg             392 non-null   float64
 1   cylinders       392 non-null   int64  
 2   displacement    392 non-null   float64
 3   horsepower      392 non-null   float64
 4   weight          392 non-null   int64  
 5   acceleration    392 non-null   float64
 6   model_year     392 non-null   int64  
 7   origin          392 non-null   object  
 8   name            392 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```

Data Visualization


```
In [100]: sns.pairplot(d,x_vars=['displacement','horsepower','weight','acceleration'],'mpg')
```

```
Out[100]: <seaborn.axisgrid.PairGrid at 0x141e5a92520>
```



```
In [30]: sns.regplot(x='displacement',y='mpg',data=d)
```

```
Out[30]: <AxesSubplot:xlabel='displacement', ylabel='mpg'>
```



Define Target variable y and x

```
In [48]: d.columns
```

```
Out[48]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',  
               'acceleration', 'model_year', 'origin', 'name'],  
              dtype='object')
```

```
In [49]: y=d['mpg']
```

```
In [50]: y.shape
```

```
Out[50]: (392,)
```

```
In [52]: x=d[['displacement','horsepower','weight','acceleration']]  
x.shape
```

```
Out[52]: (392, 4)
```

Scaling Data

```
In [56]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
x=ss.fit_transform(x)
x
```

```
Out[56]: array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
 [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
 [ 1.1825422  ,  1.18439658,  0.54038176, -1.64818924],
 ...,
 [-0.56847897, -0.53247413, -0.80463202, -1.4304305  ],
 [-0.7120053  , -0.66254009, -0.41562716,  1.11008813],
 [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])
```

```
In [103]: pd.DataFrame(x).describe()
```

```
Out[103]:
```

	0	1	2	3
count	3.920000e+02	3.920000e+02	3.920000e+02	3.920000e+02
mean	-2.537653e-16	-4.392745e-16	5.607759e-17	6.117555e-16
std	1.001278e+00	1.001278e+00	1.001278e+00	1.001278e+00
min	-1.209563e+00	-1.520975e+00	-1.608575e+00	-2.736983e+00
25%	-8.555316e-01	-7.665929e-01	-8.868535e-01	-6.410551e-01
50%	-4.153842e-01	-2.853488e-01	-2.052109e-01	-1.499869e-02
75%	7.782764e-01	5.600800e-01	7.510927e-01	5.384714e-01
max	2.493416e+00	3.265452e+00	2.549061e+00	3.360262e+00

TrainTest split data

```
In [105]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=25)
```

```
In [62]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[62]: ((274, 4), (118, 4), (274,), (118,))
```

Linear Regression Model

```
In [101]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[101]: LinearRegression()
```

```
In [106]: lr.intercept_
```

```
Out[106]: 23.6889216106858
```

```
In [83]: lr.coef_
```

```
Out[83]: array([-0.13510042, -1.4297211 , -5.23891463,  0.22436094])
```

Equation:

Mileage=23.4-0.13 Displacement-1.42 Horsepower-5.28 weight-0.224Acceleration+error

Predict Test Data

```
In [79]: y_pred=lr.predict(x_test)
```

```
In [86]: y_pred
```

```
Out[86]: array([25.24954801, 26.85525431, 26.58882904, 29.48052754, 23.91216916,
 14.9529791 , 30.0607685 , 34.07634195, 30.550342 , 11.31024173,
 18.14067535, 18.75305197, 29.80678264, 33.19954312, 17.23635872,
 16.06983768, 25.94812038, 21.15777548, 29.92508087, 25.05587641,
 22.85575427, 30.96630956, 22.82202336, 24.04513247, 25.95102384,
 26.21136844, 14.91805111, 31.85928917, 21.95227216, 26.85446824,
  8.94214825, 26.21244694, 30.20552304,  7.15733458, 26.31771126,
 30.54356872, 14.13603243, 31.02810818, 33.19140036, 31.74995879,
 11.07428823, 30.50398808, 29.36195486, 31.022648 , 23.53384962,
 22.87821543, 11.03531446, 14.3757476 , 31.44484893, 26.64255441,
 27.96470623, 21.80486111, 20.32272978, 31.27632871, 24.83127389,
 19.13391479, 28.2786737 , 25.21468804, 26.89045676, 28.76603057,
 19.03600671, 29.49310219, 28.42147856, 26.6112997 ,  7.384747 ,
 20.13152225, 22.77931428, 20.50765035, 32.81875326, 27.92430623,
 13.34341223,  8.03767139, 25.34229398, 17.23635872, 33.03710336,
 31.07878627, 21.58700058, 24.53266643, 30.38829664, 17.84737111,
 31.30622407, 30.1021144 , 22.81248978, 20.01904445,  9.12644754,
 24.50457451, 29.57695629, 29.45235437, 31.59169567, 26.49442535,
 30.32795983, 12.36145993, 16.48933189, 15.27329229, 32.77989962,
 27.25863029, 11.07878871, 25.72147567, 12.57968624, 30.4363069 ,
 27.56306784, 24.92600083, 16.21791725, 23.89776551, 18.63499966,
 10.21748386, 21.60970196, 23.01257072, 27.30850629, 30.45961552,
 29.43254102, 27.21176721, 24.2365775 , 28.87030773, 21.16703179,
 27.97152628, 24.54560958, 32.23487944])
```

Model Accuracy

```
In [91]: from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
mean_absolute_error(y_test, y_pred)
```

```
Out[91]: 3.417654680078562
```

```
In [88]: mean_absolute_percentage_error(y_test, y_pred)
```

```
Out[88]: 0.16282215595698366
```

```
In [92]: r2_score(y_test, y_pred)
```

```
Out[92]: 0.6767436309121446
```

Polynomial Regression

```
In [107]: from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=2, interaction_only=True, include_bias=False)
```

```
In [108]: x_train2 = poly.fit_transform(x_train)
```

```
In [109]: x_test2 = poly.fit_transform(x_test)
```

```
In [110]: lr.fit(x_train2, y_train)
```

```
Out[110]: LinearRegression()
```

```
In [111]: lr.intercept_
```

```
Out[111]: 21.45712035519168
```

```
In [112]: lr.coef_
```

```
Out[112]: array([-1.97594907e+00, -5.50639326e+00, -1.82341405e+00, -8.04049934e-01,
                1.55534517e+00, -4.40583099e-01, -5.33735335e-01, 1.29466895e+00,
                2.61553723e-03, 5.86761939e-01])
```

```
In [114]: y_pred_poly = lr.predict(x_test2)
```

Model Accuracy

```
In [115]: from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score
mean_absolute_error(y_test, y_pred_poly)
```

```
Out[115]: 2.924007242447458
```

```
In [116]: mean_absolute_percentage_error(y_test,y_pred_poly)
```

```
Out[116]: 0.12874881331071994
```

```
In [117]: r2_score(y_test,y_pred_poly)
```

```
Out[117]: 0.7198303534964863
```

Hand Written Digit Classification


```
In [118]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

importing data from sklearn library

```
In [119]: from sklearn.datasets import load_digits
```

```
In [120]: df=load_digits()
```

```
In [126]: _,axes=plt.subplots(nrows=1,ncols=4,figsize=(10,3))
for ax,image,label in zip(axes,df.images,df.target):
    ax.set_axis_off()
    ax.imshow(image,cmap=plt.cm.gray_r,interpolation='nearest')
    ax.set_title("Training: %i" % label)
```



Flatten image

```
In [128]: df.images.shape
```

```
Out[128]: (1797, 8, 8)
```



```
In [131]: df.images[0]
```

```
Out[131]: array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
 [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
 [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
 [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
 [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
 [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
 [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
 [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

```
In [132]: df.images[0].shape
```

```
Out[132]: (8, 8)
```

```
In [133]: n_samples=len(df.images)
data=df.images.reshape((n_samples,-1))
```

```
In [134]: data[0]
```

```
Out[134]: array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
 15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
 12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
  0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
 10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

```
In [135]: data[0].shape
```

```
Out[135]: (64,)
```

```
In [136]: data.shape
```

```
Out[136]: (1797, 64)
```

Scaling data

```
In [137]: data.min()
```

```
Out[137]: 0.0
```

```
In [138]: data.max()
```

```
Out[138]: 16.0
```

```
In [139]: data=data/16
```

```
In [140]: data.shape
```

```
Out[140]: (1797, 64)
```

Train Test split Data

```
In [141]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(data,df.target,train_size=0.7,rand
```

```
In [142]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[142]: ((1257, 64), (540, 64), (1257,), (540,))
```

Random forest model

```
In [148]: from sklearn.ensemble import RandomForestClassifier  
rf=RandomForestClassifier()
```

```
In [149]: rf.fit(x_train,y_train)
```

```
Out[149]: RandomForestClassifier()
```

Predict test data

```
In [150]: y_pred=rf.predict(x_test)
y_pred
```

```
Out[150]: array([9, 5, 9, 6, 9, 8, 3, 1, 4, 0, 1, 2, 9, 1, 7, 1, 5, 6, 8, 6, 7, 5,
 5, 2, 1, 4, 0, 5, 7, 8, 4, 2, 1, 2, 0, 4, 2, 1, 9, 9, 2, 7, 5, 3,
 9, 2, 4, 7, 8, 4, 0, 8, 7, 7, 8, 7, 8, 0, 7, 3, 3, 1, 9, 2, 4, 0,
 5, 2, 5, 1, 4, 1, 4, 9, 7, 3, 3, 6, 9, 6, 9, 7, 9, 2, 7, 2, 1, 5,
 2, 4, 1, 3, 4, 9, 7, 1, 8, 1, 3, 1, 5, 2, 5, 8, 3, 8, 2, 7, 1, 1,
 1, 1, 6, 2, 2, 3, 4, 1, 4, 6, 8, 5, 4, 0, 0, 2, 7, 0, 3, 2, 1, 6,
 9, 3, 5, 9, 6, 7, 1, 3, 8, 6, 7, 5, 8, 3, 9, 0, 1, 9, 7, 3, 2, 6,
 7, 4, 8, 1, 4, 0, 6, 3, 6, 6, 5, 5, 2, 8, 9, 4, 9, 5, 2, 1, 6, 6,
 7, 1, 0, 3, 0, 2, 8, 5, 3, 3, 5, 2, 1, 8, 4, 3, 5, 2, 7, 2, 9, 9,
 2, 4, 6, 0, 1, 6, 1, 5, 4, 4, 5, 7, 9, 8, 3, 9, 3, 5, 0, 2, 6, 4,
 1, 3, 1, 9, 6, 5, 7, 5, 1, 6, 4, 2, 1, 7, 2, 0, 2, 1, 9, 5, 7, 3,
 0, 3, 9, 9, 4, 0, 1, 2, 4, 5, 1, 0, 7, 5, 3, 7, 8, 4, 8, 3, 0, 3,
 8, 9, 3, 6, 3, 1, 0, 3, 8, 2, 2, 6, 5, 5, 6, 6, 5, 7, 0, 9, 8, 8,
 4, 0, 7, 6, 1, 8, 0, 1, 7, 0, 2, 4, 7, 8, 6, 8, 3, 2, 4, 5, 6, 0,
 7, 6, 1, 5, 5, 3, 4, 7, 0, 7, 0, 8, 3, 8, 2, 9, 8, 5, 2, 1, 2, 5,
 6, 5, 9, 1, 3, 8, 1, 7, 8, 1, 7, 8, 0, 3, 7, 5, 5, 4, 7, 3, 8, 6,
 9, 5, 7, 7, 5, 6, 5, 3, 1, 2, 7, 4, 4, 6, 5, 5, 9, 5, 2, 7, 2, 5,
 1, 9, 7, 9, 6, 8, 0, 4, 6, 9, 9, 6, 8, 2, 7, 6, 7, 4, 7, 5, 5, 6,
 5, 2, 3, 3, 0, 3, 4, 9, 3, 4, 3, 9, 2, 4, 2, 0, 5, 0, 1, 6, 8, 4,
 4, 2, 9, 0, 8, 4, 6, 9, 0, 9, 6, 0, 5, 1, 5, 5, 2, 5, 7, 9, 1, 3,
 6, 5, 8, 4, 3, 7, 5, 8, 1, 8, 2, 3, 0, 0, 0, 3, 1, 9, 0, 0, 2, 4,
 5, 0, 4, 0, 6, 0, 1, 9, 4, 8, 3, 8, 5, 9, 0, 5, 0, 5, 0, 6, 3, 5,
 2, 9, 7, 4, 6, 5, 0, 9, 2, 0, 8, 2, 9, 0, 0, 1, 0, 4, 3, 3, 9, 8,
 1, 1, 6, 3, 8, 8, 5, 9, 5, 4, 1, 7, 1, 3, 4, 5, 1, 4, 3, 2, 9, 7,
 4, 9, 0, 2, 3, 8, 1, 5, 6, 3, 7, 1])
```

```
In [151]: from sklearn.metrics import confusion_matrix,classification_report
```

```
In [152]: confusion_matrix(y_test,y_pred)
```

```
Out[152]: array([[52,  0,  0,  0,  0,  0,  0,  0,  0,  0],
 [ 0, 58,  0,  0,  0,  0,  0,  0,  0,  0],
 [ 1,  1, 54,  0,  0,  0,  0,  0,  0,  0],
 [ 0,  0,  0, 52,  0,  0,  0,  0,  0,  0],
 [ 0,  0,  0,  0, 51,  0,  0,  1,  1,  0],
 [ 0,  0,  0,  1,  0, 65,  1,  0,  0,  3],
 [ 0,  0,  0,  0,  0,  1, 46,  0,  1,  0],
 [ 0,  0,  0,  0,  0,  0,  0, 50,  0,  1],
 [ 0,  1,  0,  1,  0,  0,  0,  0, 46,  0],
 [ 0,  0,  0,  2,  0,  0,  0,  1,  0, 49]], dtype=int64)
```

```
In [153]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.98	1.00	0.99	52
1	0.97	1.00	0.98	58
2	1.00	0.96	0.98	56
3	0.93	1.00	0.96	52
4	1.00	0.96	0.98	53
5	0.98	0.93	0.96	70
6	0.98	0.96	0.97	48
7	0.96	0.98	0.97	51
8	0.96	0.96	0.96	48
9	0.92	0.94	0.93	52
accuracy			0.97	540
macro avg	0.97	0.97	0.97	540
weighted avg	0.97	0.97	0.97	540

```
In [ ]:
```